

Conversion Rate Analysis

Howard Nguyen

```
library(dplyr) # Data Wrangling and Manipulation
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

Load data

```
conversionsDF <- read.csv(file="bank-additional-full.csv", header=TRUE, sep=";")
```

```
# Shape of conversionsDF
```

```
dim(conversionsDF)
```

```
## [1] 41188    21
```

```
# Quick look at conversionsDF
```

```
head(conversionsDF)
```

```
##   age      job marital  education default housing loan  contact month
## 1  56 housemaid married  basic.4y      no      no  no telephone  may
## 2  57 services married high.school unknown      no  no  no telephone  may
## 3  37 services married high.school      no      yes  no telephone  may
## 4  40 admin. married  basic.6y      no      no  no telephone  may
## 5  56 services married high.school      no      no  yes telephone  may
## 6  45 services married  basic.9y unknown      no  no  no telephone  may
##  day_of_week duration campaign pdays previous  poutcome emp.var.rate
## 1      mon      261         1    999         0 nonexistent         1.1
## 2      mon      149         1    999         0 nonexistent         1.1
```

```
## 3      mon      226      1  999      0 nonexistent      1.1
## 4      mon      151      1  999      0 nonexistent      1.1
## 5      mon      307      1  999      0 nonexistent      1.1
## 6      mon      198      1  999      0 nonexistent      1.1
##   cons.price.idx cons.conf.idx euribor3m nr.employed y
## 1      93.994      -36.4      4.857      5191 no
## 2      93.994      -36.4      4.857      5191 no
## 3      93.994      -36.4      4.857      5191 no
## 4      93.994      -36.4      4.857      5191 no
## 5      93.994      -36.4      4.857      5191 no
## 6      93.994      -36.4      4.857      5191 no
```

the output variable, y, which has information on whether a client has subscribed to a term deposit, is encoded as 'yes' or 'no'. In order to simplify our conversion rate computations, we will encode this variable as 1 for 'yes' and 0 for 'no'.

```
# Encode conversions as 0s and 1s
conversionsDF$y <- ifelse(conversionsDF$y=="yes",1,0)
conversionsDF$conversion <- as.integer(conversionsDF$y)
```

```
tail(conversionsDF)
```

```
##      age      job marital      education default housing loan  contact
## 41183  29 unemployed single      basic.4y      no      yes  no cellular
## 41184  73      retired married professional.course      no      yes  no cellular
## 41185  46 blue-collar married professional.course      no      no  no cellular
## 41186  56      retired married university.degree      no      yes  no cellular
## 41187  44 technician married professional.course      no      no  no cellular
## 41188  74      retired married professional.course      no      yes  no cellular
##      month day_of_week duration campaign pdays previous      poutcome
## 41183  nov      fri      112      1      9      1      success
## 41184  nov      fri      334      1  999      0 nonexistent
## 41185  nov      fri      383      1  999      0 nonexistent
## 41186  nov      fri      189      2  999      0 nonexistent
## 41187  nov      fri      442      1  999      0 nonexistent
## 41188  nov      fri      239      3  999      1      failure
##      emp.var.rate cons.price.idx cons.conf.idx euribor3m nr.employed y
## 41183      -1.1      94.767      -50.8      1.028      4963.6 0
## 41184      -1.1      94.767      -50.8      1.028      4963.6 1
## 41185      -1.1      94.767      -50.8      1.028      4963.6 0
## 41186      -1.1      94.767      -50.8      1.028      4963.6 0
## 41187      -1.1      94.767      -50.8      1.028      4963.6 1
## 41188      -1.1      94.767      -50.8      1.028      4963.6 0
##      conversion
## 41183      0
## 41184      1
## 41185      0
## 41186      0
## 41187      1
## 41188      0
```

1. Aggregate Conversion Rate

Since we have already encoded the output variable as 1 for those who have converted and 0 for those who have not in a column, named conversion, we can simply sum over this column to get the total number of conversions. The following code snippet shows how we can sum over the conversion column and get the total number of clients in the data:

```
# total of numbers of conversions
sprintf("total conversions: %i out of %i", sum(conversionsDF$conversion), nrow(conversionsDF))

## [1] "total conversions: 4640 out of 41188"

# total number of clients in the data (= number of records in the data)
sprintf("conversion rate: %0.2f%%", sum(conversionsDF$conversion)/nrow(conversionsDF)*100.0)

## [1] "conversion rate: 11.27%"
```

2. Conversion Rates by Age

Aggregate conversion rate tells us the overall performance of our marketing campaign. However, it does not give us that much insight. When we are reporting and tracking the progress of marketing efforts, we typically would want to dive deeper into the data and break down the customer base into multiple segments and compute KPIs for individual segments. We will first break our data into smaller segments by age and see how the conversion rates differ by different age groups.

Note: The pipe operator, `%>%`, in this code, is the way you can apply different functions sequentially. In this code snippet, we are passing `conversionsDF` to a `group_by` function, then passing the results of this `group_by` function to the `summarise` function, and lastly to the `mutate` function.

In the `group_by` function, we are grouping the `DataFrame` by the column `age`. Then, for each age group, we are counting the number of records in each group, by using a function, `n()`, and naming it `TotalCount`. Also, we are summing over the column, `conversion`, for each age group, by using the `sum` function, and naming it `NumConversions`.

Lastly, we are using the `mutate` function, which adds new variables, while preserving the original `DataFrame`, to compute conversion rates for each age group. As we can see, we are simply dividing `NumConversion` by `TotalCount` and multiplying it by 100 to get the conversion rates.

```
# a. by age
conversionsByAge <- conversionsDF %>%
  group_by(Age=age) %>%
  summarise(TotalCount=n(), NumConversions=sum(conversion)) %>%
  mutate(ConversionRate=NumConversions/TotalCount*100.0)

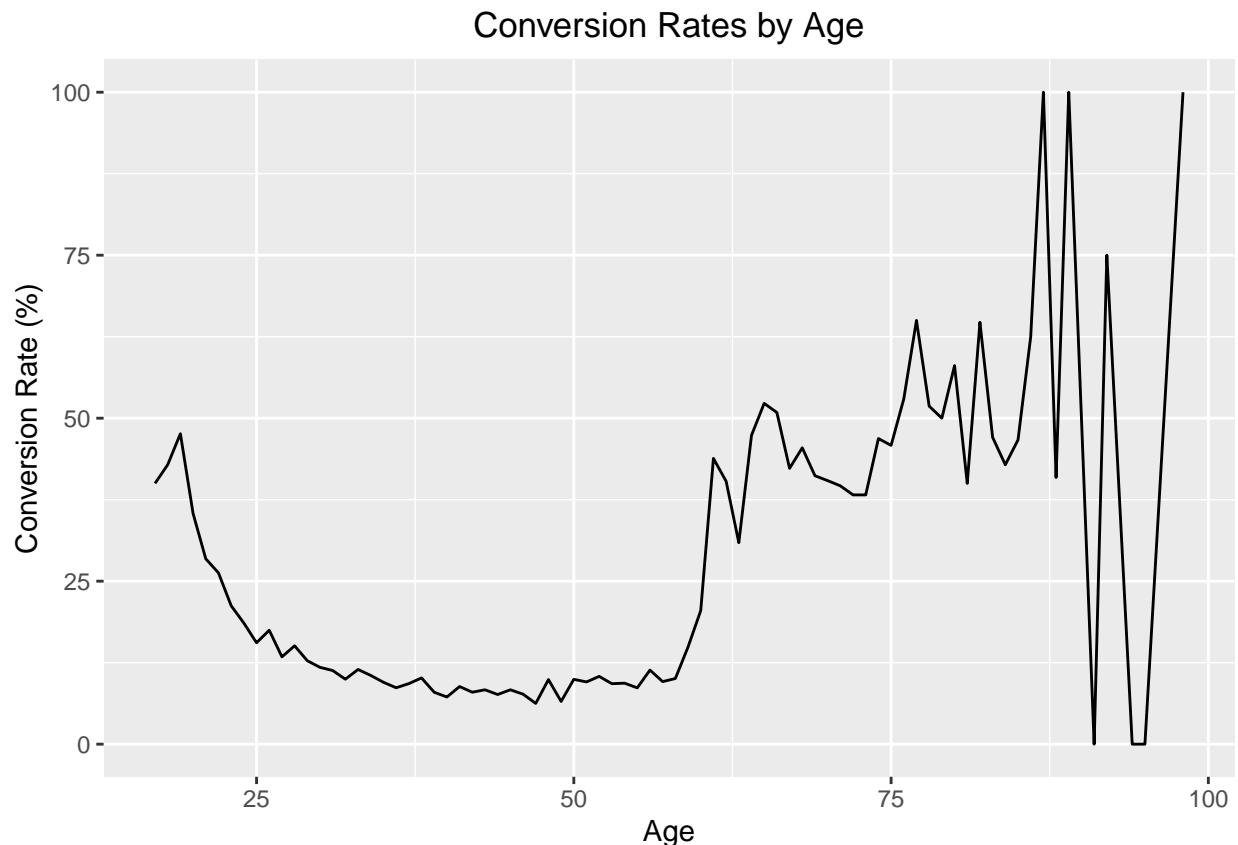
head(conversionsByAge)

## # A tibble: 6 x 4
##   Age TotalCount NumConversions ConversionRate
##   <int>      <int>      <int>      <dbl>
## 1    17         5           2         40
## 2    18        28          12        42.9
## 3    19        42          20        47.6
## 4    20        65          23        35.4
## 5    21       102          29        28.4
## 6    22       137          36        26.3
```

Another way to look at conversion rates across client ages is by plotting a line chart, as shown in the following screenshot:

Note: We are using the `ggplot` function to initialize a `ggplot` object with `conversionsByAge` as the data and the column, `Age`, as the x-axis and the column, `ConversionRate`, as the y-axis. Then, we use `geom_line` function to connect the observations and create a line chart. You can change the title of a plot, by using `ggtitle` function. Also, you can use `xlab` and `ylab` functions to rename the x-axis label and y-axis label respectively.

```
# line chart
ggplot(data=conversionsByAge, aes(x=Age, y=ConversionRate)) +
  geom_line() +
  ggtitle('Conversion Rates by Age') +
  xlab("Age") +
  ylab("Conversion Rate (%)") +
  theme(plot.title = element_text(hjust = 0.5))
```



One thing that is noticeable in the previous line chart is the fact that there seems to be lots of noise in older age groups. Conversion rates for those who are 70 or older vary a lot and if we look at the data, this is mostly because the number of clients in this age group is relatively small, compared to other age groups.

In order to reduce this unwanted noise, we can group multiple ages together. In the section below, we group bank clients into six different groups, based on their age—between 18 and 30, between 30 and 40, between 40 and 50, between 50 and 60, between 60 and 70, and 70 and older. The following code can be used to group the clients into their corresponding groups:

```
# b. by age groups
conversionsByAgeGroup <- conversionsDF %>%
  group_by(AgeGroup=cut(age, breaks= seq(20, 70, by = 10)) ) %>%
  summarise(TotalCount=n(), NumConversions=sum(conversion)) %>%
  mutate(ConversionRate=NumConversions/TotalCount*100.0)

conversionsByAgeGroup$AgeGroup <- as.character(conversionsByAgeGroup$AgeGroup)
conversionsByAgeGroup$AgeGroup[6] <- "70+"
```

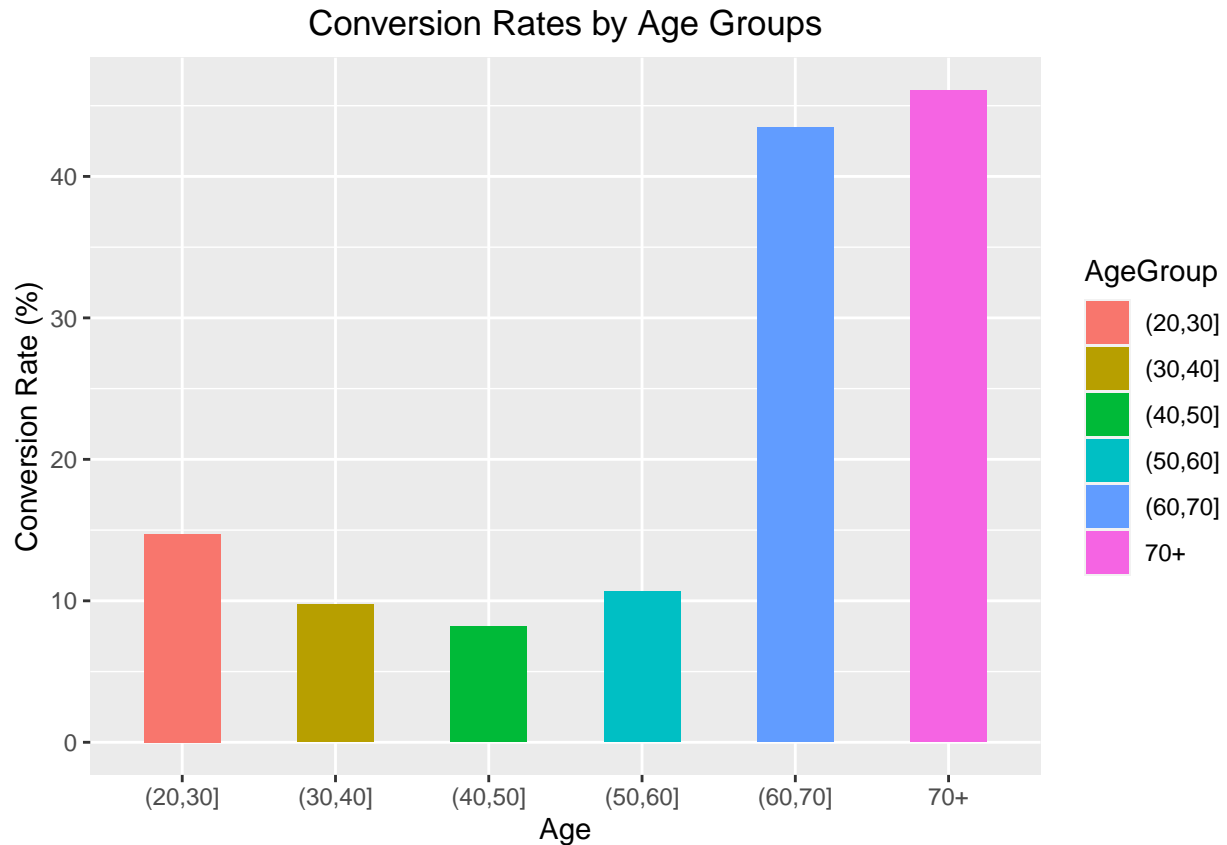
```
conversionsByAgeGroup
```

```
## # A tibble: 6 x 4
##   AgeGroup TotalCount NumConversions ConversionRate
##   <chr>         <int>         <int>         <dbl>
## 1 (20,30]       7243           1067          14.7
## 2 (30,40]      16385           1597           9.75
## 3 (40,50]      10240            837           8.17
## 4 (50,60]       6270            668          10.7
## 5 (60,70]       488             212          43.4
## 6 70+           562             259          46.1
```

As with the previous case, we are using the `group_by` function to group the `conversionsDF` data by the age column. The difference here is how we used the `cut` function to create the age range for each age group.

The `breaks` argument defines the points at which the `cut` function is going to divide the `DataFrame`. The argument, `seq(20, 70, by = 10)`, means we are going to create a sequence from 20 to 70 in increments of 10. Once the data is grouped by these age groups, the rest are the same as before. We are using the `summarise` and `mutate` functions to compute for the `TotalCount`, `NumConversions`, and `ConversionRate` columns.

```
# bar chart
ggplot(conversionsByAgeGroup, aes(x=AgeGroup, y=ConversionRate, fill=AgeGroup)) +
  geom_bar(width=0.5, stat="identity") +
  ggtitle('Conversion Rates by Age Groups') +
  xlab("Age") +
  ylab("Conversion Rate (%)") +
  theme(plot.title = element_text(hjust = 0.5))
```



3. Conversion Rates by Number of Contacts

```

conversionsByNumContact <- conversionsDF %>%
  group_by(NumContact=campaign) %>%
  summarise(TotalCount=n(), NumConversions=sum(conversion)) %>%
  mutate(ConversionRate=NumConversions/TotalCount*100.0)

head(conversionsByNumContact, 10)

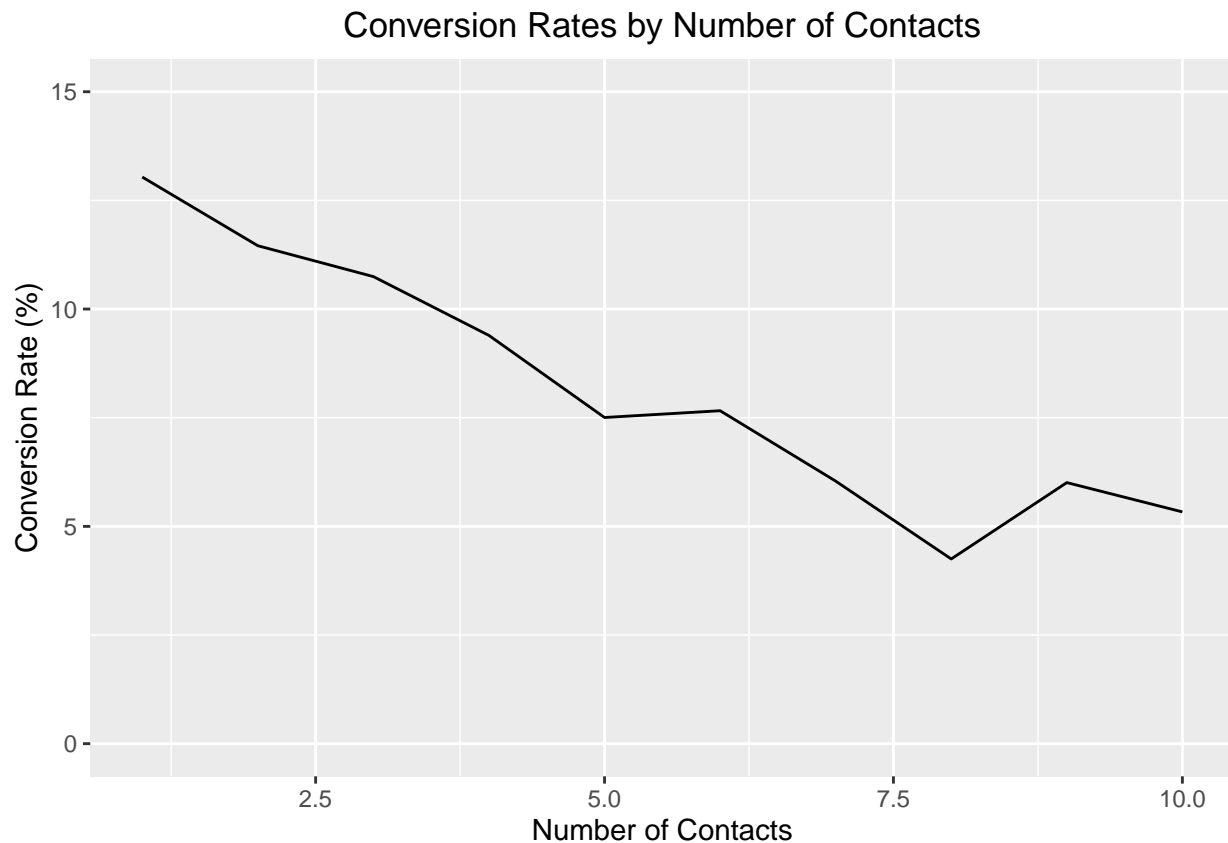
```

```

## # A tibble: 10 x 4
##   NumContact TotalCount NumConversions ConversionRate
##   <int>      <int>      <int>      <dbl>
## 1         1    17642         2300         13.0
## 2         2    10570         1211         11.5
## 3         3     5341          574         10.7
## 4         4     2651          249          9.39
## 5         5     1599          120          7.50
## 6         6       979           75          7.66
## 7         7       629           38          6.04
## 8         8       400           17          4.25
## 9         9       283           17          6.01
## 10        10       225           12          5.33

```

```
# line chart
ggplot(data=head(conversionsByNumContact, 10), aes(x=NumContact, y=ConversionRate)) +
  geom_line() +
  ggtitle('Conversion Rates by Number of Contacts') +
  xlab("Number of Contacts") +
  ylab("Conversion Rate (%)") +
  ylim(c(0, 15)) +
  theme(plot.title = element_text(hjust = 0.5))
```



4. Conversions vs. Non-Conversions

One other thing we can look at is the demographic differences between the converted clients and non-converted clients. This type of analysis can help us identify what differentiates converted groups from non-converted groups in our marketing campaigns and helps us understand our target clients better and what types of customers respond better to our marketing efforts. In this exercise, we will compare the distributions of the marital status among the conversions and non-conversions groups.

We will first count the number of conversions and non-conversions for each marital status. The following code shows how we can compute this using R functions:

```
# 4.1. Marital Status
conversionsByMaritalStatus <- conversionsDF %>%
  group_by(Marital=marital, Conversion=conversion) %>%
  summarise(Count=n())
```

'summarise()' has grouped output by 'Marital'. You can override using the '.groups' argument.

```
conversionsByMaritalStatus
```

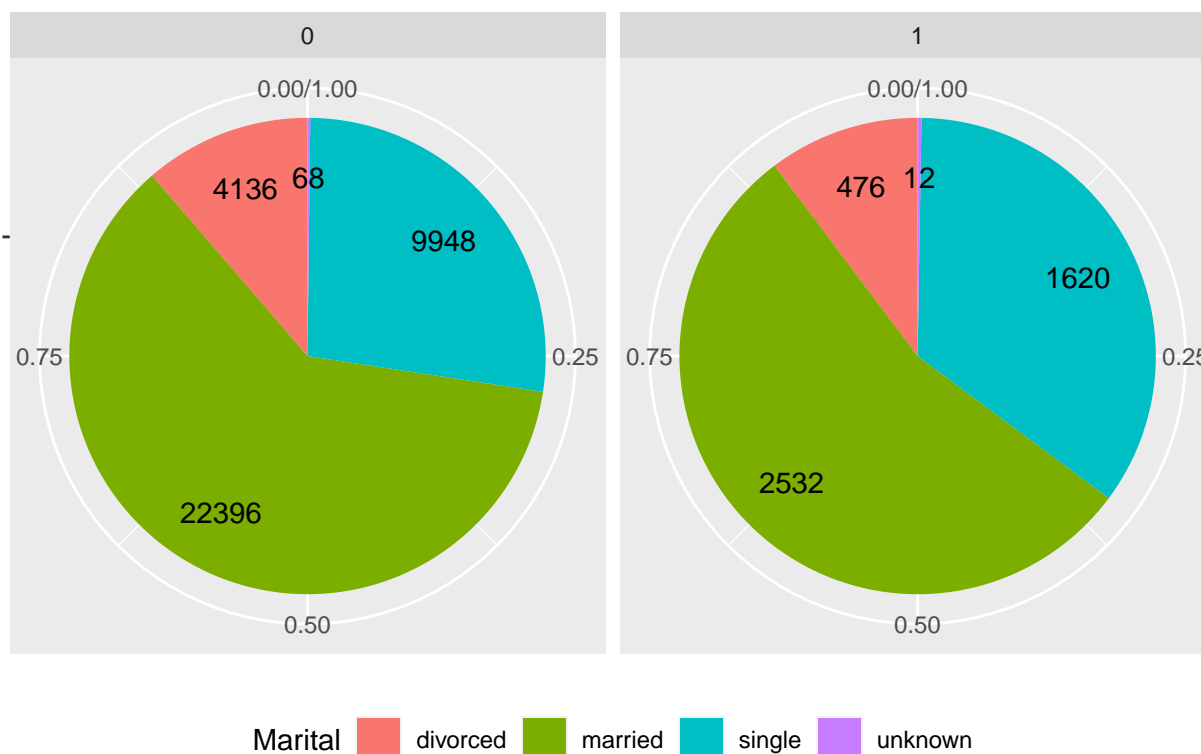
```
## # A tibble: 8 x 3
## # Groups:   Marital [4]
##   Marital Conversion Count
##   <chr>      <int> <int>
## 1 divorced      0  4136
## 2 divorced      1   476
## 3 married       0 22396
## 4 married       1  2532
## 5 single        0  9948
## 6 single        1  1620
## 7 unknown       0    68
## 8 unknown       1    12
```

Note: As we can see from the above code, we are using the pipe operator, %>%, in the dplyr package to pass the DataFrame, conversionsDF to the group_by function and then to the summarise function. In the group_by function, we are grouping by two columns, marital and conversion. In the summarise function, we are simply counting the number of records in each group, by using the n function.

Another way to present the above data by pie-chart:

```
# pie chart
ggplot(conversionsByMaritalStatus, aes(x="", y=Count, fill=Marital)) +
  geom_bar(width=1, stat = "identity", position=position_fill()) +
  geom_text(aes(x=1.25, label=Count), position=position_fill(vjust = 0.5)) +
  coord_polar("y") +
  facet_wrap(~Conversion) +
  ggtitle('Marital Status (0: Non Conversions, 1: Conversions)') +
  theme(
    axis.title.x=element_blank(),
    axis.title.y=element_blank(),
    plot.title=element_text(hjust=0.5),
    legend.position='bottom'
  )
```


Marital Status (0: Non Conversions, 1: Conversions)



Note: For building a pie chart in R, we are using the same `geom_bar` function, just as if we are building a bar chart. The difference here is `coord_polar("y")`, which transforms a bar chart into a pie chart. Then, we are using the `facet_wrap` function to create two columns of pie charts by the column, `Conversion`. This builds two pie charts, one for the conversions group and another for the non-conversions group.

Compared to the tabular format of the data output, pie charts make it much easier to understand the overall distributions of the data. With pie charts, we can easily see that the married group takes up the largest proportions in both conversions and non-conversions groups, while the single group comes second. Using pie charts, we can easily visualize the similarities and differences between two groups.

4.2. Education

```
conversionsByEducation <- conversionsDF %>%
  group_by(Education=education, Conversion=conversion) %>%
  summarise(Count=n())
```

'summarise()' has grouped output by 'Education'. You can override using the '.groups' argument.

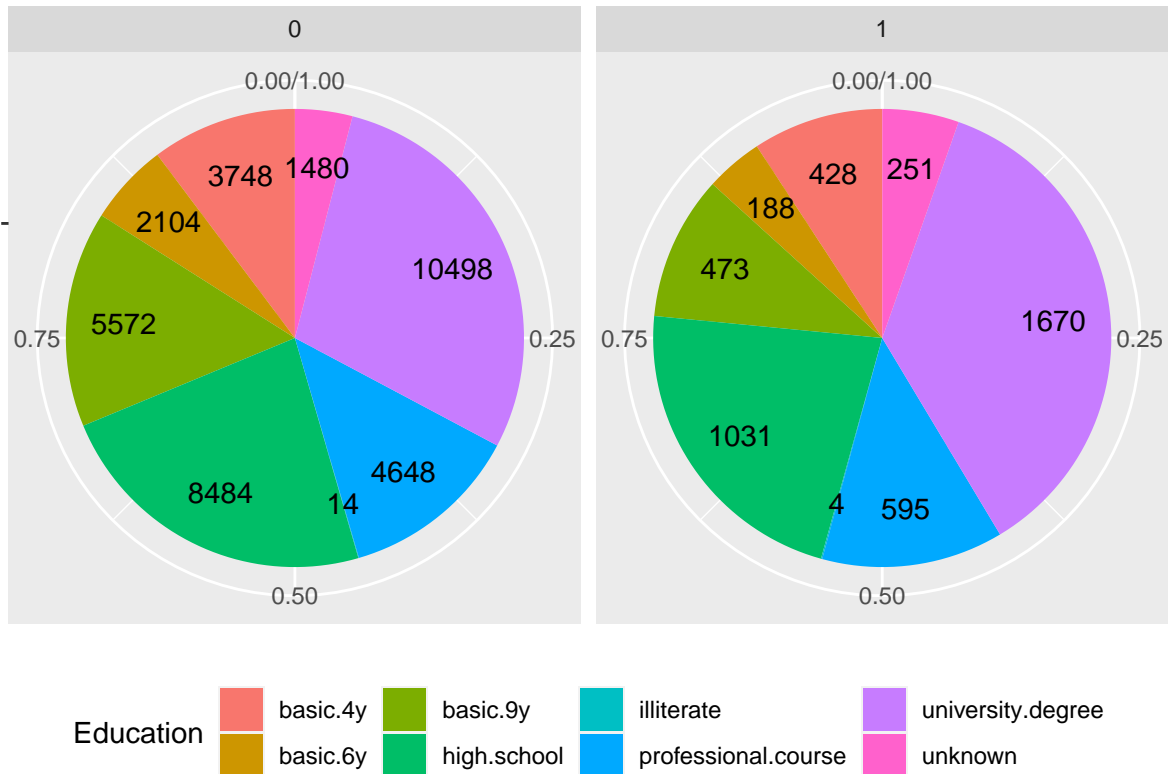
```
conversionsByEducation
```

```
## # A tibble: 16 x 3
## # Groups:   Education [8]
##   Education      Conversion Count
##   <chr>          <int> <int>
## 1 basic.4y             0  3748
## 2 basic.4y             1   428
## 3 basic.6y             0  2104
```

##	4	basic.6y	1	188
##	5	basic.9y	0	5572
##	6	basic.9y	1	473
##	7	high.school	0	8484
##	8	high.school	1	1031
##	9	illiterate	0	14
##	10	illiterate	1	4
##	11	professional.course	0	4648
##	12	professional.course	1	595
##	13	university.degree	0	10498
##	14	university.degree	1	1670
##	15	unknown	0	1480
##	16	unknown	1	251

```
# pie chart
ggplot(conversionsByEducation, aes(x="", y=Count, fill=Education)) +
  geom_bar(width=1, stat = "identity", position=position_fill()) +
  geom_text(aes(x=1.25, label=Count), position=position_fill(vjust = 0.5)) +
  coord_polar("y") +
  facet_wrap(~Conversion) +
  ggtitle('Education (0: Non Conversions, 1: Conversions)') +
  theme(
    axis.title.x=element_blank(),
    axis.title.y=element_blank(),
    plot.title=element_text(hjust=0.5),
    legend.position='bottom'
  )
```

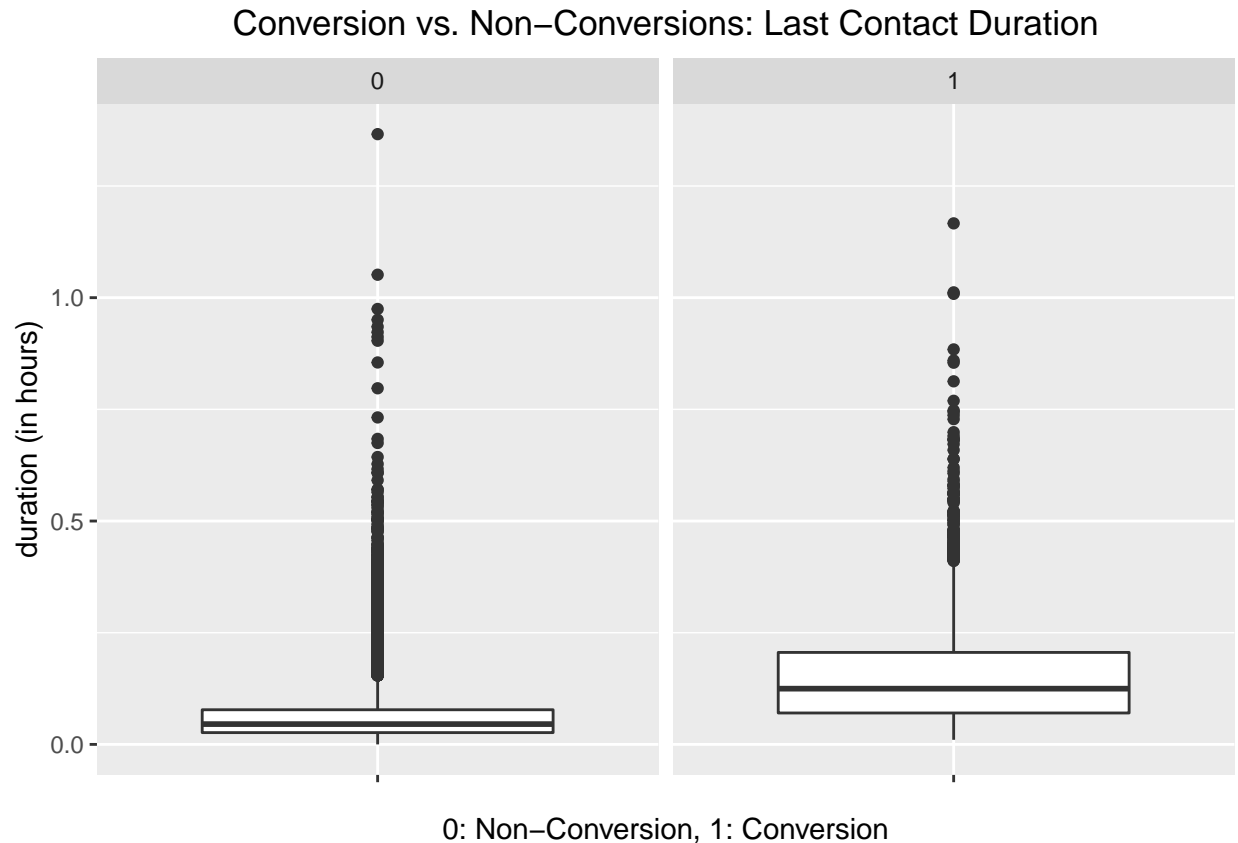
Education (0: Non Conversions, 1: Conversions)



4.3. Last Contact Duration

```
conversionsDF$duration <- conversionsDF$duration / (60*60)
```

```
ggplot(conversionsDF, aes(x="", y=duration)) +
  geom_boxplot() +
  facet_wrap(~conversion) +
  ylab("duration (in hours)") +
  xlab("0: Non-Conversion, 1: Conversion") +
  ggtitle("Conversion vs. Non-Conversions: Last Contact Duration") +
  theme(plot.title=element_text(hjust=0.5))
```



5. Conversions by Age Groups & Marital Status

So far, we have aggregated our data by one criterion. However, there are cases where you want to group the data by more than one column. In this section, we will discuss how we can analyze and report conversion rates by more than one criterion. In this section, we will use the age groups that we have built in the previous section and the marital status as the two columns to group by.

Let's first look at the code:

```
conversionsByAgeMarital <- conversionsDF %>%
  group_by(AgeGroup=cut(age, breaks= seq(20, 70, by = 10)), Marital=marital) %>%
  summarise(Count=n(), NumConversions=sum(conversion)) %>%
  mutate(TotalCount=sum(Count)) %>%
  mutate(ConversionRate=NumConversions/TotalCount)
```

'summarise()' has grouped output by 'AgeGroup'. You can override using the '.groups' argument.

```
conversionsByAgeMarital$AgeGroup <- as.character(conversionsByAgeMarital$AgeGroup)
conversionsByAgeMarital$AgeGroup[is.na(conversionsByAgeMarital$AgeGroup)] <- "70+"
conversionsByAgeMarital
```

```
## # A tibble: 23 x 6
## # Groups:   AgeGroup [6]
##   AgeGroup Marital Count NumConversions TotalCount ConversionRate
```

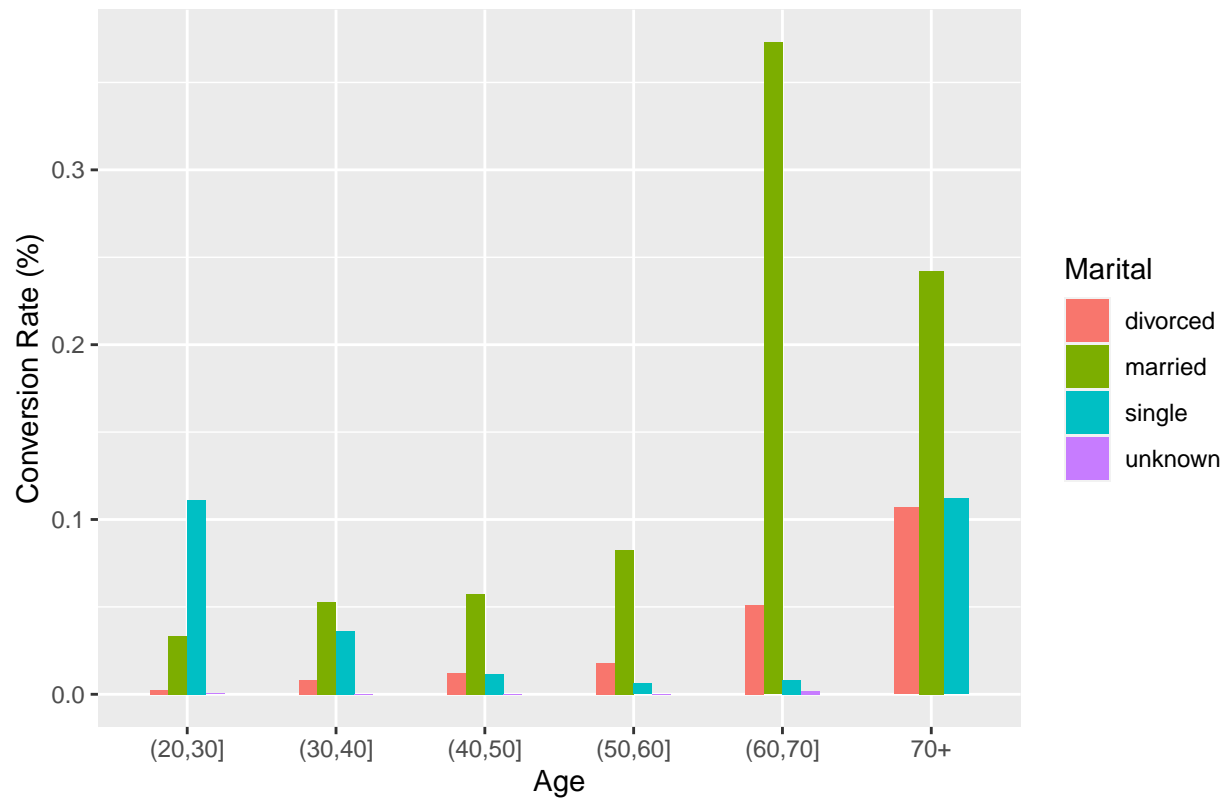
```
##      <chr>      <chr>      <int>          <int>          <int>          <dbl>
## 1 (20,30] divorced    229             18           7243      0.00249
## 2 (20,30] married    2389            242           7243      0.0334
## 3 (20,30] single     4612            804           7243      0.111
## 4 (20,30] unknown     13              3           7243      0.000414
## 5 (30,40] divorced    1505            135          16385      0.00824
## 6 (30,40] married    9705            867          16385      0.0529
## 7 (30,40] single     5139            591          16385      0.0361
## 8 (30,40] unknown     36              4           16385      0.000244
## 9 (40,50] divorced    1548            126          10240      0.0123
## 10 (40,50] married    7383            588          10240      0.0574
## # ... with 13 more rows
```

Note: Similar to when we built custom age groups, we are using the cut function in group_by to create age groups from 20 to 70 in increments of 10. However, “we are grouping by the column,marital, as well this time.


Then, we are using the summarise function to compute the number of records in each group Count, and the number of conversions in each group, NumConversions. Then, using the mutate function, we calculate the total counts in each age group, named TotalCount, and the conversion rates in each group, named ConversionRate.

```
# bar chart
ggplot(conversionsByAgeMarital, aes(x=AgeGroup, y=ConversionRate, fill=Marital)) +
  geom_bar(width=0.5, stat="identity", position="dodge") +
  ylab("Conversion Rate (%)") +
  xlab("Age") +
  ggtitle("Conversion Rates by Age and Marital Status") +
  theme(plot.title=element_text(hjust=0.5))
```

Conversion Rates by Age and Marital Status



```
# stacked bar chart
ggplot(conversionsByAgeMarital, aes(x=AgeGroup, y=ConversionRate, fill=Marital)) +
  geom_bar(width=0.5, stat="identity", position="stack") +
  ylab("Conversion Rate (%)") +
  xlab("Age") +
  ggtitle("Conversion Rates by Age and Marital Status") +
  theme(plot.title=element_text(hjust=0.5))
```



Summary: These above are various KPIs that are often used in marketing to track the progress of marketing campaigns. We have learned how important it is to look at how much sales revenue each marketing strategy generates. When analyzing the sales revenue metrics, we have seen that it is important to approach it from different angles. We might want to look at not only the aggregate sales revenue, but also time-series (monthly, quarterly, or yearly) sales revenue. We might also want to look at sales attributed to each individual marketing campaigns and how much revenue each campaign generated for the company. We will go other sections with various metrics to analyze for digital marketing channels as well, such as CPA, CTR, lead ratio, and conversion rates.